

Python for data analysis

Simple linear Regression

```
In [149]: import numpy as np #import Libraries
import pandas as pd
import seaborn as sns
import sklearn as sl
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [150]: data=pd.read_excel(r'F:\dataset.xlsx') ## import dataset
```

In [151]: data

Out[151]:

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	3.5	34
6	5.3	34
7	4.2	22
8	1.5	53
9	2.5	24
10	4.1	12
11	9.7	47
12	6.7	27
13	8.5	23
14	4.9	30
15	7.8	34
16	9.3	34
17	7.9	22
18	5.5	53
19	6.7	24
20	3.5	54
21	6.7	30
22	7.8	34
23	9.6	76
24	8.1	34

Basic statistic of data

```
In [152]: print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Hours   25 non-null        float64
1   Scores  25 non-null        int64
dtypes: float64(1), int64(1)
memory usage: 464.0 bytes
None
```

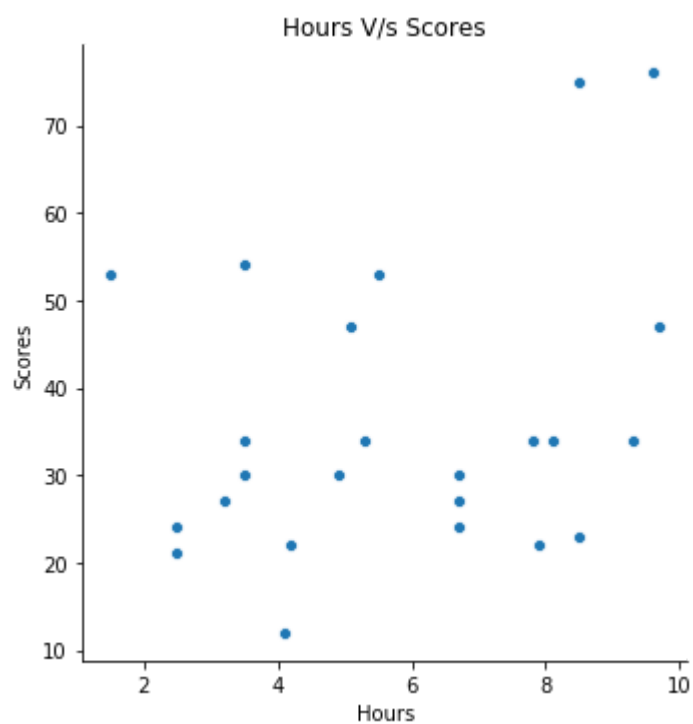
Given dataset contain 25 entries and there is no any null value present in data.

```
In [153]: print(data.describe())    ##Description of data
```

	Hours	Scores
count	25.000000	25.000000
mean	5.864000	36.040000
std	2.470641	16.084361
min	1.500000	12.000000
25%	3.500000	24.000000
50%	5.500000	34.000000
75%	7.900000	47.000000
max	9.700000	76.000000



Vizualization

```
In [154]: sns.relplot(x='Hours', y='Scores', data=data) ## vizualization of data  
plt.title("Hours V/s Scores")  
plt.show()
```



```
In [155]: sns.catplot("Hours", data=data, kind='box')  
sns.catplot("Scores", data=data, kind='box')
```

```
Out[155]: <seaborn.axisgrid.FacetGrid at 0x11c84fb0>
```




```
In [156]: cor=data.corr()    ## Correlation of data  
cor
```

```
Out[156]:
```


	Hours	Scores
Hours	1.00000	0.27779
Scores	0.27779	1.00000

```
In [157]: sns.heatmap(cor,annot=True)  
plt.show()
```



```
In [158]: sns.distplot(data['Scores'],hist=False)    ## plotting the distribution of scores of data
```

```
Out[158]: <matplotlib.axes._subplots.AxesSubplot at 0x11be85b0>
```



Preparing the data

```
In [159]: x=data.iloc[:, :-1].values  
y=data.iloc[:, 1].values
```

Train test split

```
In [160]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split( x,y,test_size=0.2,random_state  
=0 )
```

```
In [161]: from sklearn.linear_model import LinearRegression  
  
reg=LinearRegression()  
model=reg.fit(x_train,y_train)
```


```
In [162]: print(model.coef_),print(model.intercept_)  
  
[2.0095554]  
25.27543866262086
```

```
Out[162]: (None, None)
```

The best fit of line is

Scores= $12.860 + 4.577 \times \text{Hours}$

```
In [163]: line=model.intercept_+model.coef_*x  ##Plotting the regression line  
plt.scatter(x,y,color='black')  
plt.plot(x,line,color='red')  
plt.show()
```



Prediction

```
In [164]: y_pred=model.predict(x_test)  
y_pred
```

```
Out[164]: array([32.30888256, 31.70601594, 38.73945985, 43.96430389, 44.76812605])
```

```
In [165]: pd.DataFrame({'Actual':y_test,'predict':y_pred})    ##comparing Actual value versus Predict value
```

Out[165]:

	Actual	predict
0	34	32.308883
1	27	31.706016
2	24	38.739460
3	34	43.964304
4	47	44.768126

```
In [166]: Hours=10.25
Predict_Score=model.predict([[Hours]])
Predict_Score
print('No of Hours=8.5')
print("Predict Score=",format(Predict_Score[0]))
```

No of Hours=8.5
Predict Score= 45.87338151868674

```
In [167]: from sklearn import metrics    ##mean square error
print('mean absolute error:',round(metrics.mean_absolute_error(y_test,y_pred)),3)
print('mean Squared error:',round(metrics.mean_squared_error(y_test,y_pred)),3)
print('Root mean squared error:',round(np.sqrt(metrics.mean_squared_error(y_test,y_pred)),3))
```

mean absolute error: 7.0 3
mean Squared error: 69.0 3
Root mean squared error: 8.325

The value of root mean squared error is 8.325 which is greater than 10% of the mean percentage of all the students(36.04)

```
In [168]: from scipy import stats    # R- squared value
slope,inter,r,p,std=stats.linregress(y_test,y_pred)
print('R-squared Value:',r)
```

R-squared Value: 0.5724771973429205

Conclusion: Here R- Squared value is 0.5724 means the model explains 57.24% of variability in the dependent variable (Score) by independent variable (Hours).

In []: